

IS883: Deploying Generative AI

Mohannad Elhamod

Models in the wild

Model Types

We are not going to get into technical details, but certain models may be more fit for certain tasks:

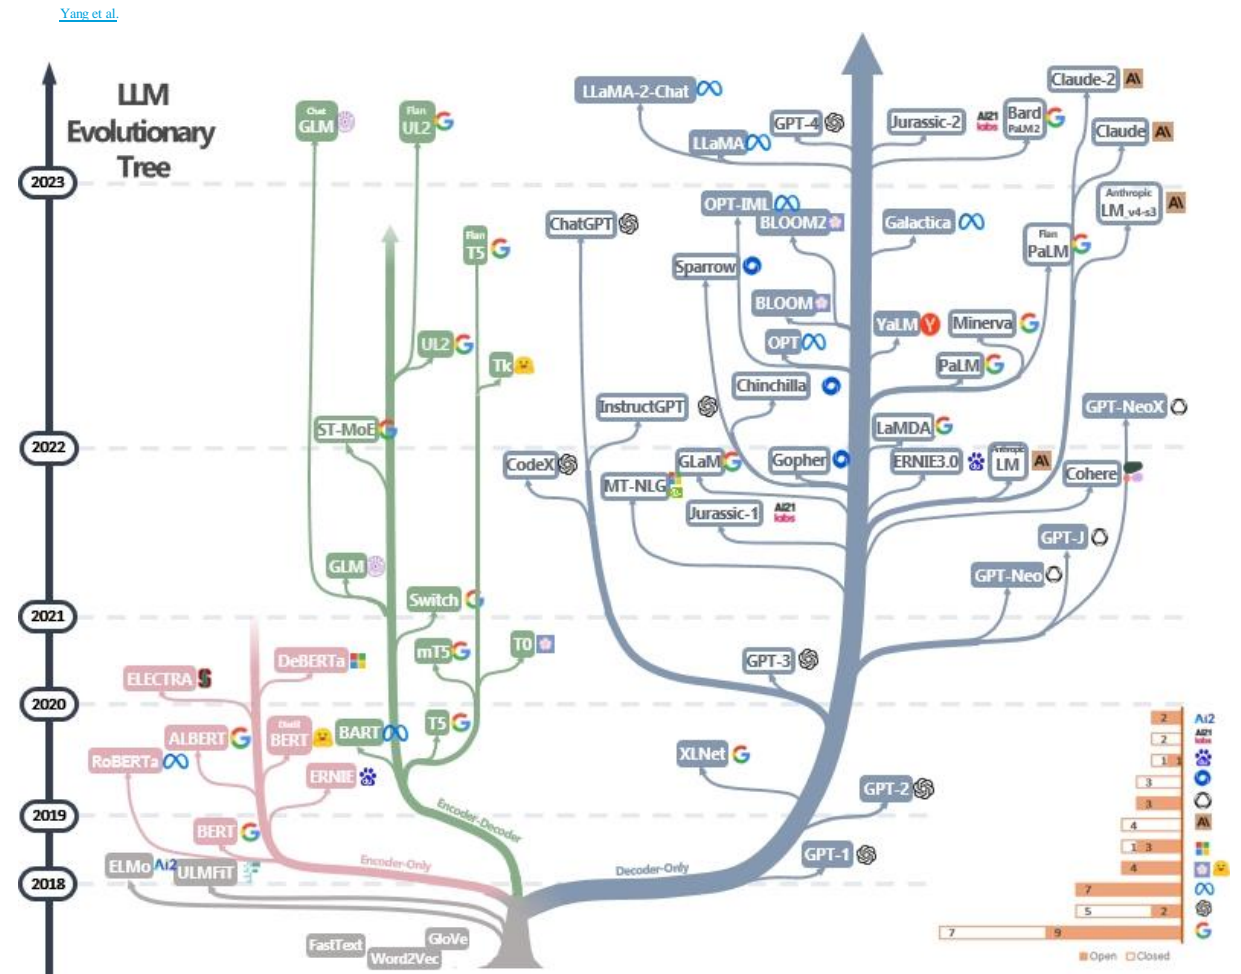
| Model | Examples | Tasks |
|-----------------|--|--|
| Encoder | ALBERT, BERT, DistilBERT, ELECTRA, RoBERTa | Sentence classification, named entity recognition, extractive question answering |
| Decoder | CTRL, GPT, GPT-2, Transformer XL | Text generation |
| Encoder-decoder | BART, T5, Marian, mBART | Summarization, translation, generative question answering |

[Lavink.org](#)

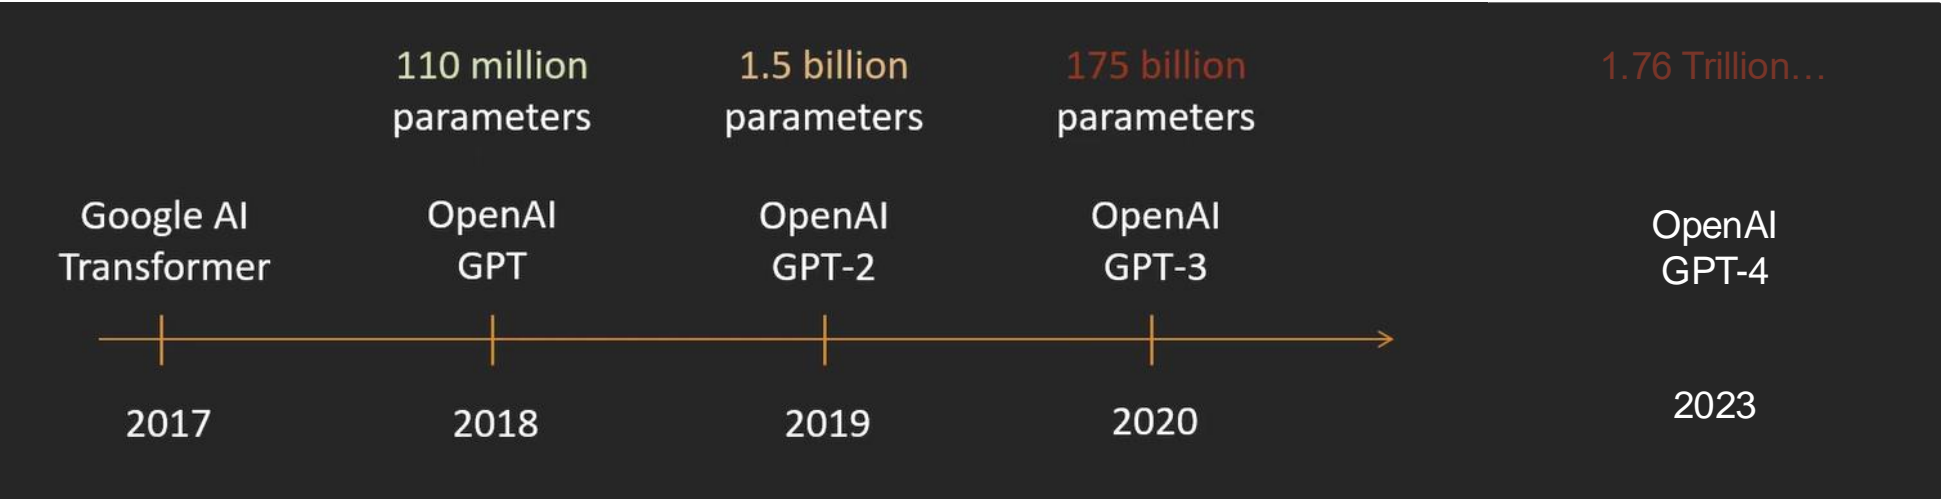

Why so many?

Where do the differences come from?

- Data.
- Model type and size.
- Hyperparameters (context size, embedding size,...).
- Training process (the cost function, fine-tuning, human feedback, etc.).



The GPT Evolution...



[AI Coffee Break with Letitia](#)

[Book Corpus](#)
[WebText](#)


1,038 books (around 74M sentences and 1G words) of 16 different sub-genres (e.g., Romance, Historical, Adventure, etc.)

[Common Crawl + ...](#)

Over 240 billion pages.
Petabytes of data.

[????](#)

The GPT Evolution...



780B tokens


Link in the description below. 📌 Chowdhery et al. 2022

Total dataset size = 780 billion tokens

| Data source | Proportion of data |
|---|--------------------|
| Social media conversations (multilingual) | 50% |
| Filtered webpages (multilingual) | 27% |
| Books (English) | 13% |
| GitHub (code) | 5% |
| Wikipedia (multilingual) | 4% |
| News (English) | 1% |

[AI Coffee Break with Letitia](#)

Different model sizes



[Jay Alumar](#)

Exploring Your Options

- [OpenAI model reference](#)
- [HuggingFace tasks](#)
- [HuggingFace models](#)

How much training does it take?

2 example models

GPT-3 (2020)

50,257 vocabulary size
2048 context length
175B parameters
Trained on 300B tokens

| Model Name | n_{params} | n_{layers} | d_{model} | n_{heads} | d_{head} | Batch Size | Learning Rate |
|-----------------------|---------------------|---------------------|--------------------|--------------------|-------------------|------------|----------------------|
| GPT-3 Small | 125M | 12 | 768 | 12 | 64 | 0.5M | 6.0×10^{-4} |
| GPT-3 Medium | 350M | 24 | 1024 | 16 | 64 | 0.5M | 3.0×10^{-4} |
| GPT-3 Large | 760M | 24 | 1536 | 16 | 96 | 0.5M | 2.5×10^{-4} |
| GPT-3 XL | 1.3B | 24 | 2048 | 24 | 128 | 1M | 2.0×10^{-4} |
| GPT-3 2.7B | 2.7B | 32 | 2560 | 32 | 80 | 1M | 1.6×10^{-4} |
| GPT-3 6.7B | 6.7B | 32 | 4096 | 32 | 128 | 2M | 1.2×10^{-4} |
| GPT-3 13B | 13.0B | 40 | 5140 | 40 | 128 | 2M | 1.0×10^{-4} |
| GPT-3 175B or "GPT-3" | 175.0B | 96 | 12288 | 96 | 128 | 3.2M | 0.6×10^{-4} |

Table 2.1: Sizes, architectures, and learning hyper-parameters (batch size in tokens and learning rate) of the models which we trained. All models were trained for a total of 300 billion tokens.

Training: (rough order of magnitude to have in mind)

- O(1,000 - 10,000) V100 GPUs
- O(1) month of training
- O(1-10) \$M

LLaMA (2023)

32,000 vocabulary size
2048 context length
65B parameters
Trained on 1-1.4T tokens

| params | dimension | n_{heads} | n_{layers} | learning rate | batch size | n_{tokens} |
|--------|-----------|--------------------|---------------------|---------------|------------|---------------------|
| 6.7B | 4096 | 32 | 32 | $3.0e^{-4}$ | 4M | 1.0T |
| 13.0B | 5120 | 40 | 40 | $3.0e^{-4}$ | 4M | 1.0T |
| 32.5B | 6656 | 52 | 60 | $1.5e^{-4}$ | 4M | 1.4T |
| 65.2B | 8192 | 64 | 80 | $1.5e^{-4}$ | 4M | 1.4T |

Table 2: Model sizes, architectures, and optimization hyper-parameters.

Training for 65B model:

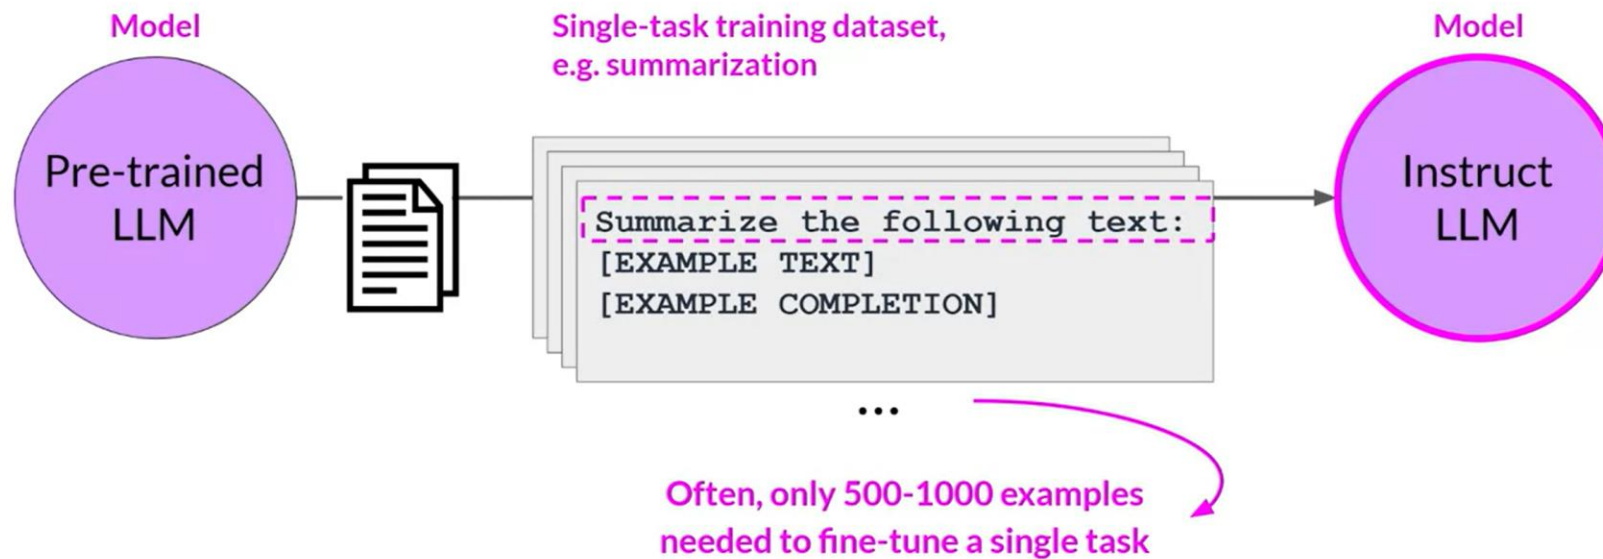
- 2,048 A100 GPUs
- 21 days of training
- \$5M

[Language Models are Few-Shot Learners, OpenAI 2020]
[LLaMA: Open and Efficient Foundation Language Models, Meta AI 2023]

Pre-trained Models: Democratizing AI

- Most of us don't have the expertise, data, or resources to train anything close to these impressive large models.
- Instead:
 - *Zero-shot Learning*: We can use open-source models out-of-the-box, even though they have never seen our data before.
 - *Transfer learning/Fine-Tuning*: Can be used as a base for further training (e.g., if the training data is non-public legal documents).

Example: Instruct LLMs



[Coursen](#)

In-Class Work

HuggingFace

Resources

- [Meaning and calculation of perplexity.](#)
- [Video: LLMs vs The Brain](#)
- [Video: Deciding which pre-trained model to fine-tune](#)