IS883: Deploying Generative AI

Mohannad Elhamod

BOSTON UNIVERSITY

Boston University Questrom School of Business

AI Truthfulness

BOSTON UNIVERSITY

Boston University Questrom School of Business

Do LLMs Know The Truth?

- Much of the training data is not factual.
  - There is a lot of misinformation out there.
  - It is based on fiction or casual conversation.

- LLMs are predicting next probably word.
  - They do not do a database lookup!
  - They have no understanding of cause and effect

Solution 1: Using Agents

- Get evidence *online* (e.g., Google Search).

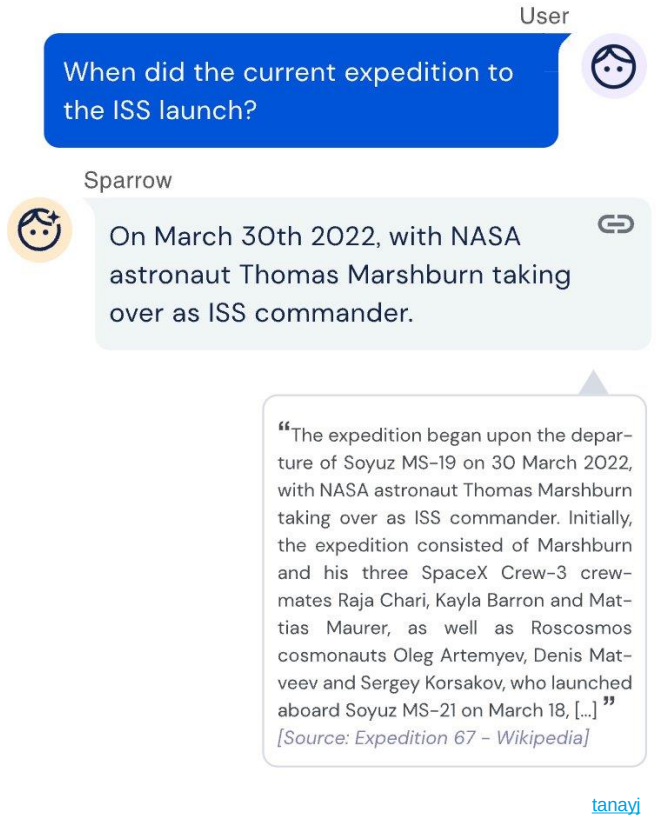

Solution 1: Using Agents

- Get evidence *offline* (e.g., a document)…
- This is called **RAG**.

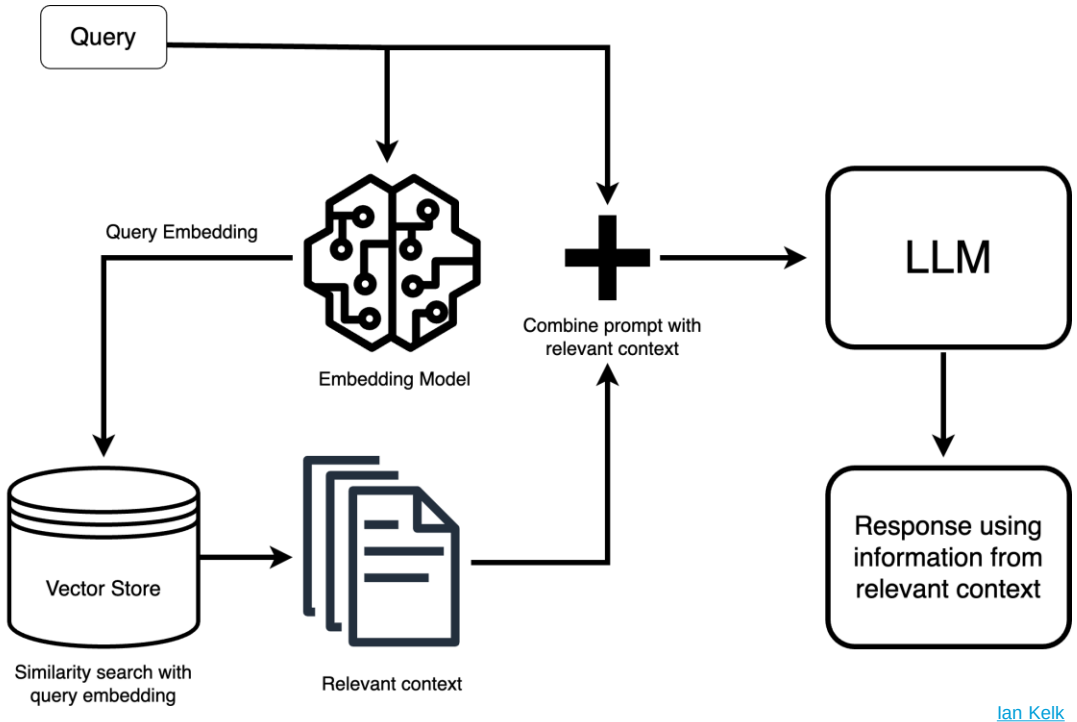

Ian Kelk

Solution 1: Using Agents

- Get evidence *offline* (e.g., a document)…
- This is called **RAG**.

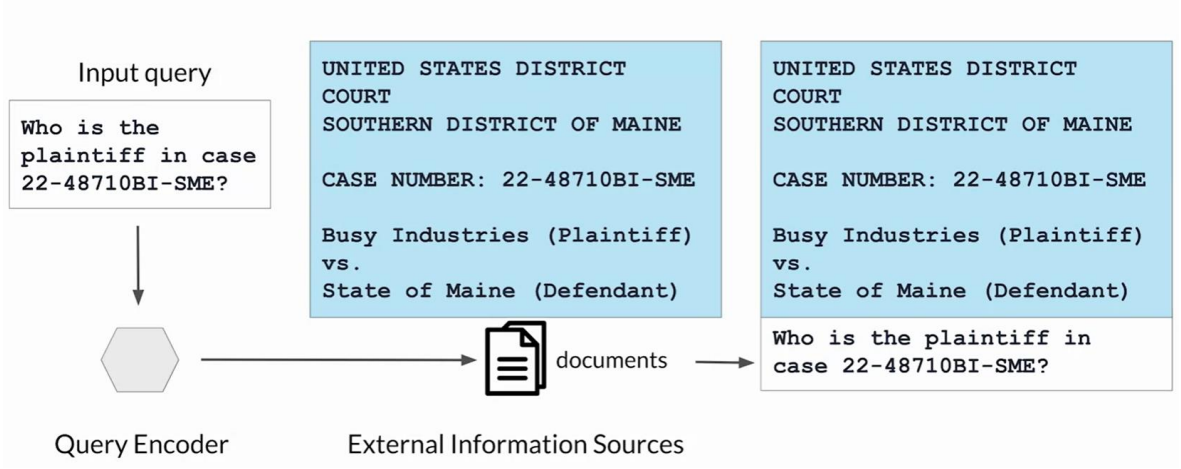

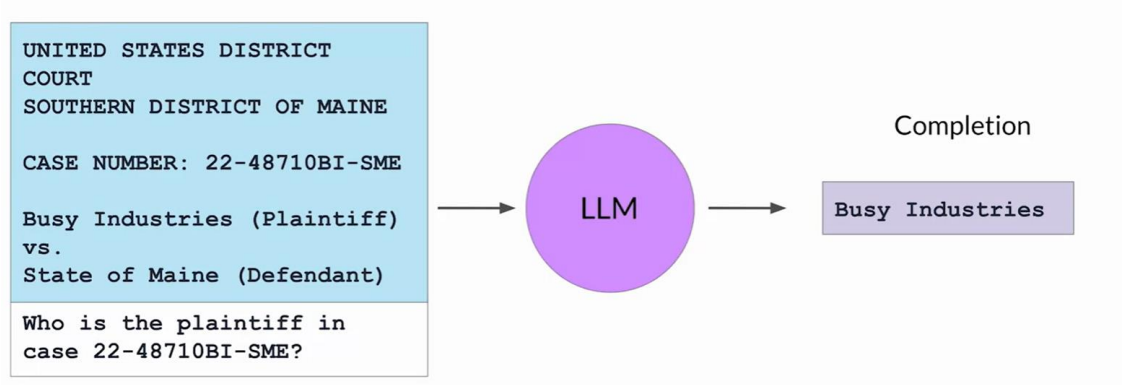

Solution 1: Using Agents

- Demo!

Solution 2: Prompt Engineering

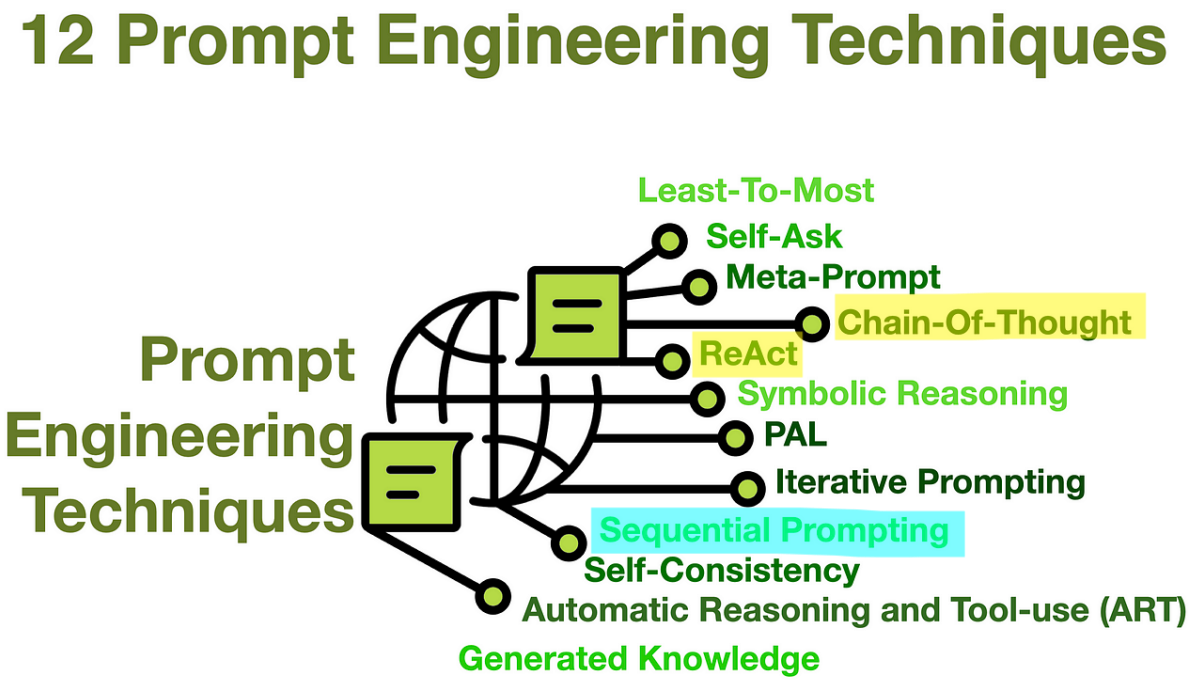

2.1. Chain of Thought (CoT)

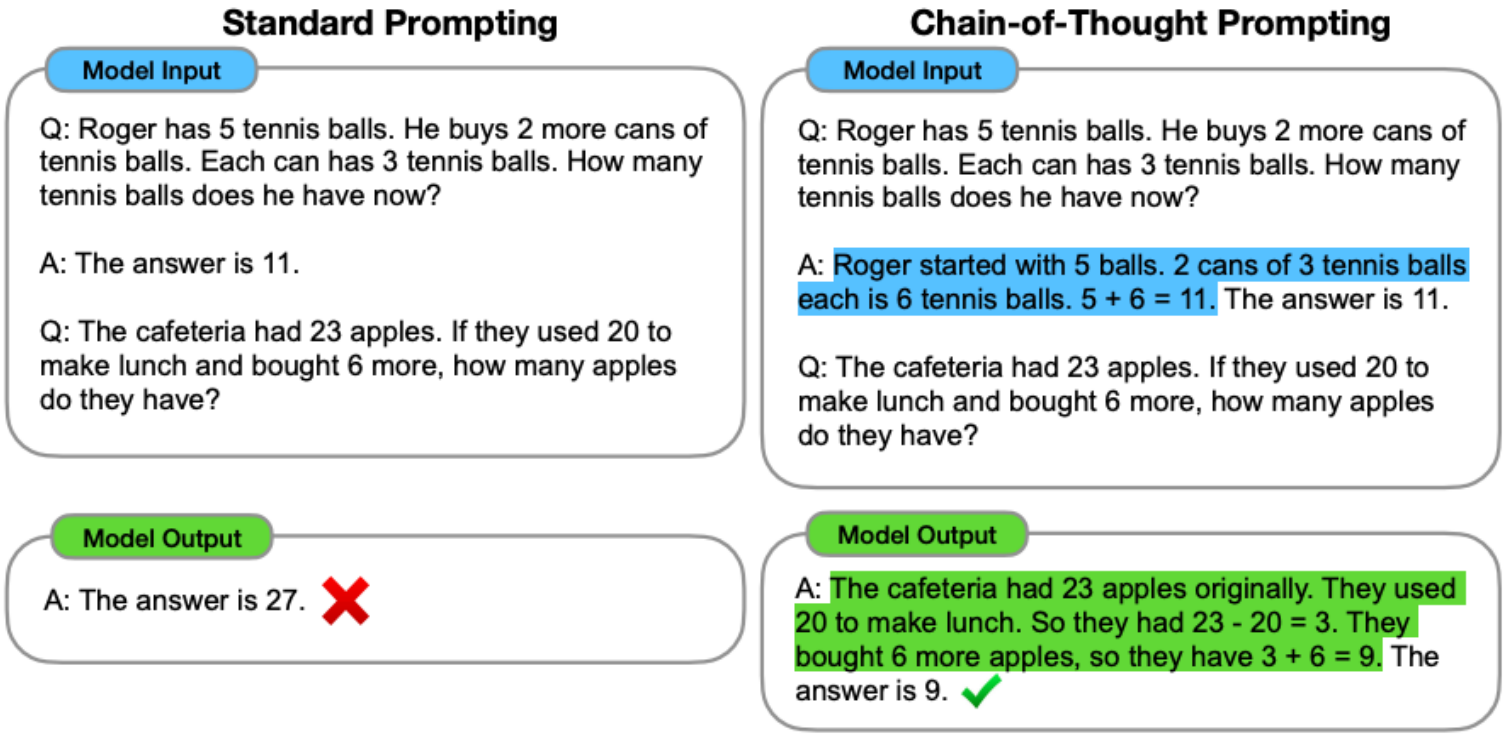

2.1. Chain of Thought (CoT)

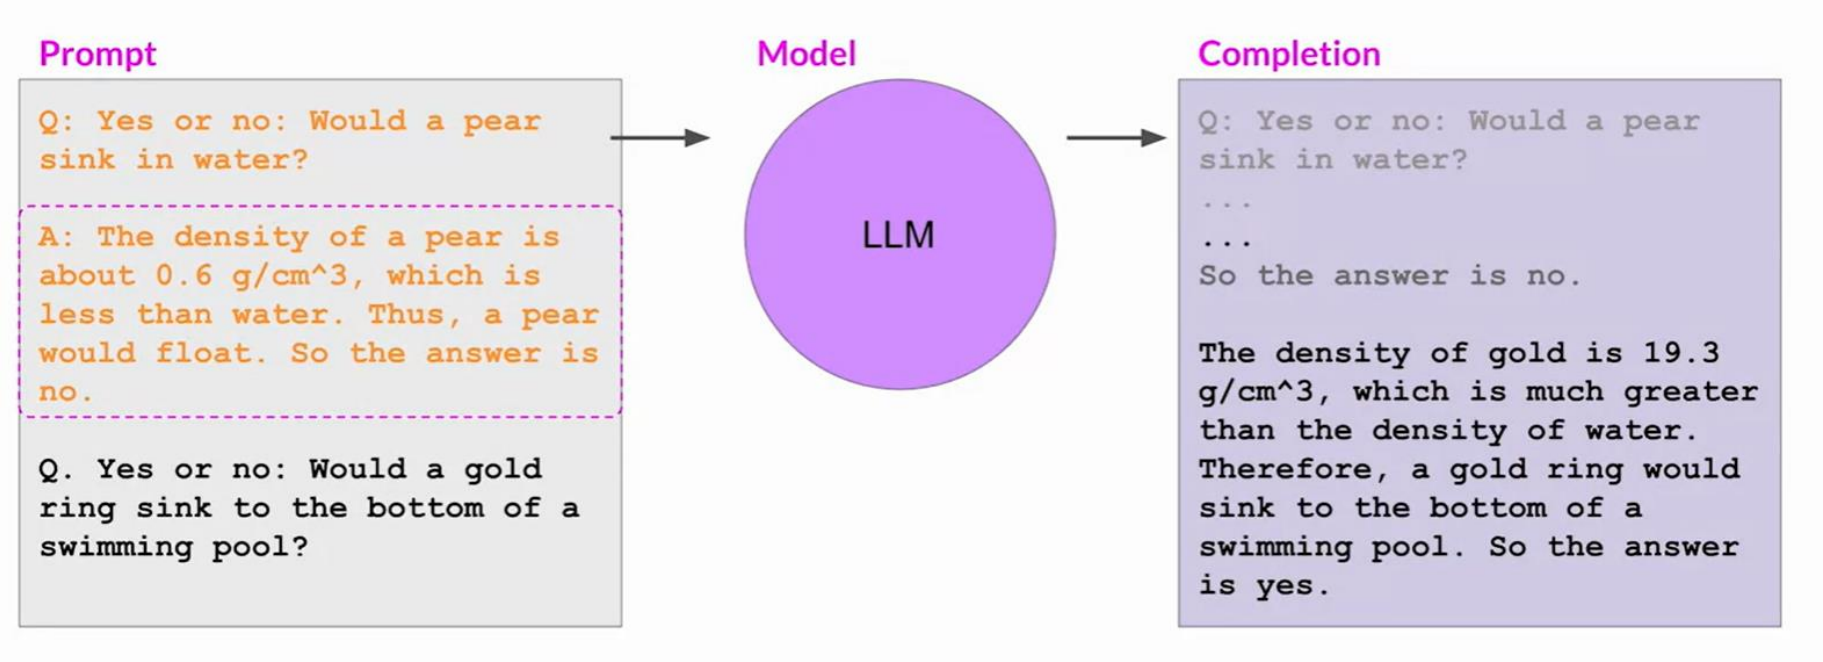

coursera.org

2.2. ReAct.

- Reasoning and Acting.

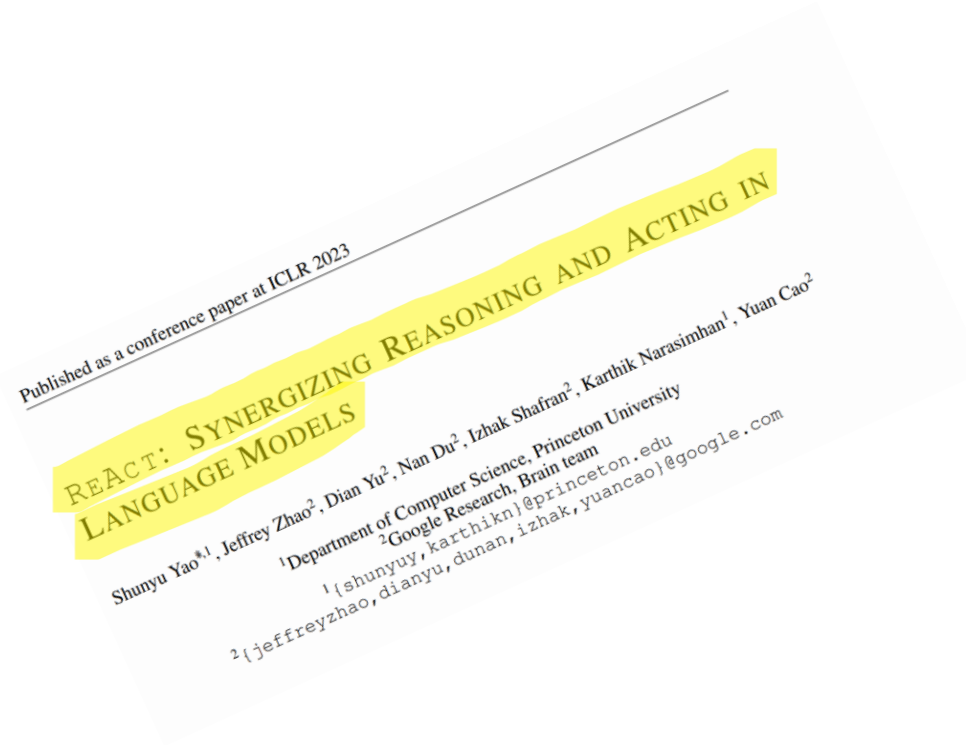

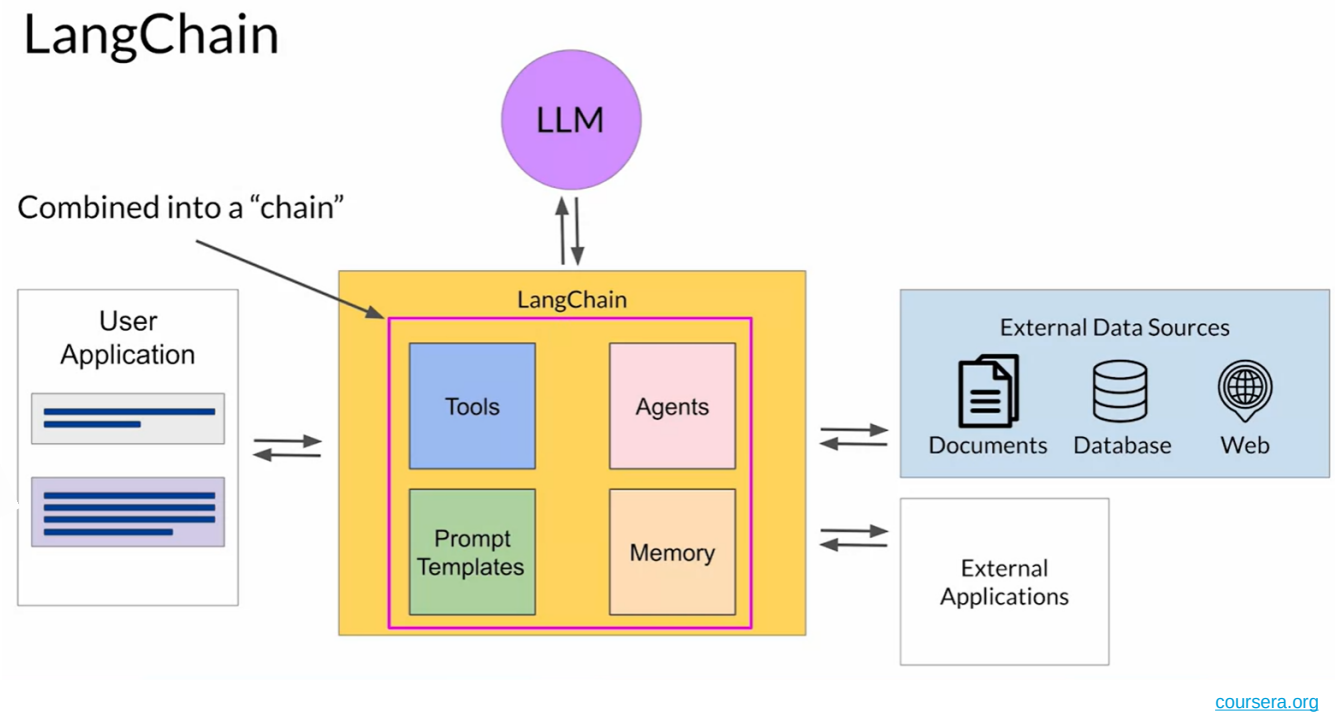

coursera.org

2.2. ReAct.

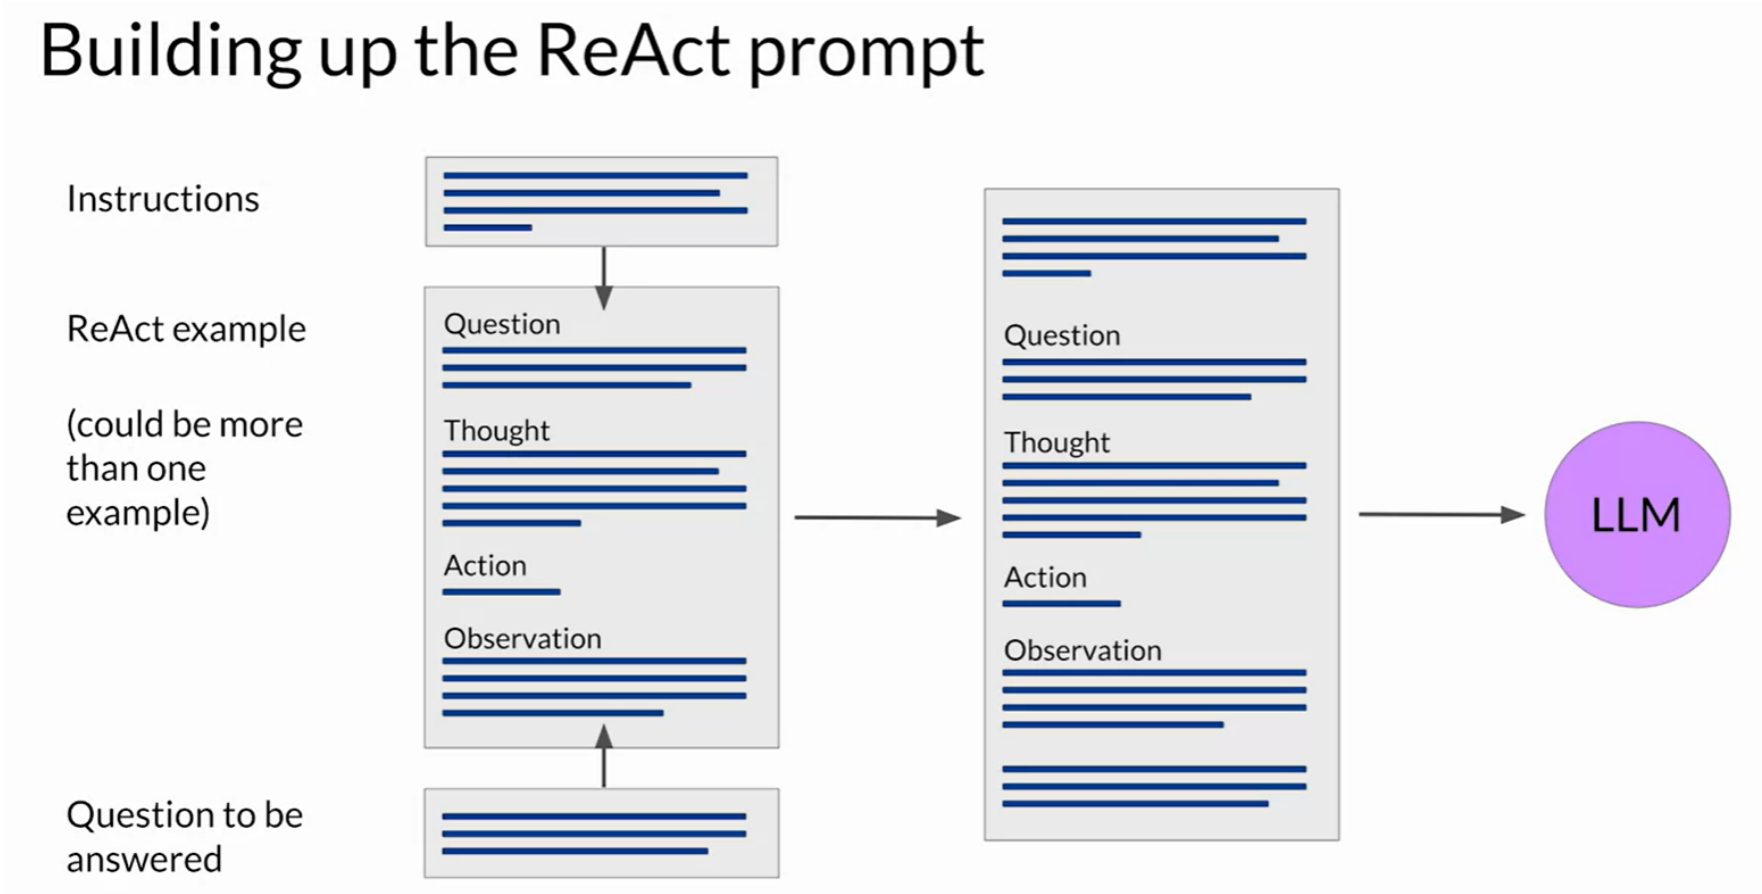

coursera.org

2.2. ReAct

LangChain output parsing works with prompt templates

```
EXAMPLES = ["""
Question: What is the elevation range
for the area that the eastern sector
of the Colorado orogeny extends into?

Thought: need to search Colorado orogeny, find
the area that the eastern sector of the Colorado
orogeny extends into, then find the elevation range
of the area.

Action: Search[Colorado orogeny]

Observation: The Colorado orogeny was an
episode of mountain building (an orogeny) in
Colorado and surrounding areas.

Thought: It does not mention the eastern sector.
So I need to look up eastern sector.
Action: Lookup[eastern sector]

...

Thought: High Plains rise in elevation from
around 1,800 to 7,000 ft, so the answer is 1,800 to
7,000 ft.

Action: Finish[1,800 to 7,000 ft]""",
]
```

LangChain library functions parse the LLM's output assuming that it will use certain keywords.

Example here uses **Thought**, **Action**, **Observation** as keywords for Chain-of-Thought Reasoning. (ReAct)

deeplearning.ai

2.3 Insertion of Misinformation

BOSTON UNIVERSITY

Insertion of Misinformation

– *False info prompt* (FIP): The prompt includes false information related to the question. For example:

> ✗ False Information: "Alfred Hitchcock directed 2001: A Space Odyssey."
> Question: "Who directed 2001: A Space Odyssey?"
>
> ---
>
> ✓ Correct Answer: "Stanley Kubrick"

– *Random info prompt* (RIP): The prompt includes random, unrelated information. For example:

> ✱ Random Information: "In the 1960s, video recorders were first developed."
> Question: "Who directed 2001: A Space Odyssey?"
>
> ---
>
> ✓ Correct Answer: "Stanley Kubrick"

A. Fastowski et al.

Insertion of Misinformation

Prompt V1:
⋮ ⋮ ⋮ ⋮
Respond with the exact answer only.

Prompt V2:
⋮ ⋮ ⋮ ⋮
Respond with the true, exact answer only.

- Using TriviaQA dataset

	GPT-4o		GPT-3.5		Mistral-7B		LLaMA-2-13B	
	Prompt V1	Prompt V2	Prompt V1	Prompt V2	Prompt V1	Prompt V2	Prompt V1	Prompt V2
B	0.987	0.986	0.982	0.971	1.000	0.984	0.829	0.815
RIP	0.958	0.940	0.914	0.908	0.866	0.846	0.734	0.706
FIP	0.921	0.934	0.781	0.863	0.516	0.539	0.359	0.364
FIP×2	0.759	0.853	0.642	0.739	0.352	0.376	0.231	0.269
FIP×5	0.710	0.820	0.592	0.678	0.287	0.304	0.182	0.203
FIP×10	0.687	0.810	0.578	0.671	0.265	0.301	0.158	0.177
% FIP×10 vs. B	-30.4%	-17.8%	-41.1%	-30.9%	-73.5%	-69.4%	-80.9%	-78.3%

A. Fastowski et al.